

시활용

21860019 서주완

Linear-Regression

보험료 예측 모델

Linear-Regression

데이터셋 분석

Age – 나이

Sex – 성별

Bmi – 체지방지수

Children – 자녀수

Smoker – 흡연여부

Region – 거주지역

Expenses - 보험료

Linear-Regression

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from sklearn.model_selection import train_test_split
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

df = pd.read_csv('/content/drive/MyDrive/ai/insurance.csv')
```

Linear-Regression

```
df.head(5)  
df.isnull().sum().sort_values(ascending=False)
```

```
age      0  
sex      0  
bmi      0  
children 0  
smoker   0  
region   0  
expenses 0  
dtype: int64
```

df

	age	sex	bmi	children	smoker	region	expenses
0	19	female	27.9	0	yes	southwest	16884.92
1	18	male	33.8	1	no	southeast	1725.55
2	28	male	33.0	3	no	southeast	4449.46
3	33	male	22.7	0	no	northwest	21984.47
4	32	male	28.9	0	no	northwest	3866.86
...
1333	50	male	31.0	3	no	northwest	10600.55
1334	18	female	31.9	0	no	northeast	2205.98
1335	18	female	36.9	0	no	southeast	1629.83
1336	21	female	25.8	0	no	southwest	2007.95
1337	61	female	29.1	0	yes	northwest	29141.36

1338 rows × 7 columns

Linear-Regression

```
le = LabelEncoder()  
df['sex'] = le.fit_transform(df['sex'])  
  
# 전처리된 데이터 확인  
print(df['sex'])
```

```
0    0  
1    1  
2    1  
3    1  
4    1
```

```
1333  1  
1334  0  
1335  0  
1336  0  
1337  0
```

Name: sex, Length: 1338, dtype: int64

```
df['smoker'] = le.fit_transform(df['smoker'])  
df['region'] = le.fit_transform(df['region'])  
  
# 전처리된 데이터 확인  
print(df['smoker'])
```

```
0    1  
1    0  
2    0  
3    0  
4    0
```

```
1333  0  
1334  0  
1335  0  
1336  0  
1337  1
```

Name: smoker, Length: 1338, dtype: int64

```
# 전처리된 데이터 확인  
print(df['region'])
```

```
0    3  
1    2  
2    2  
3    1  
4    1
```

```
1333  1  
1334  0  
1335  2  
1336  3  
1337  1
```

Name: region, Length: 1338, dtype: int64

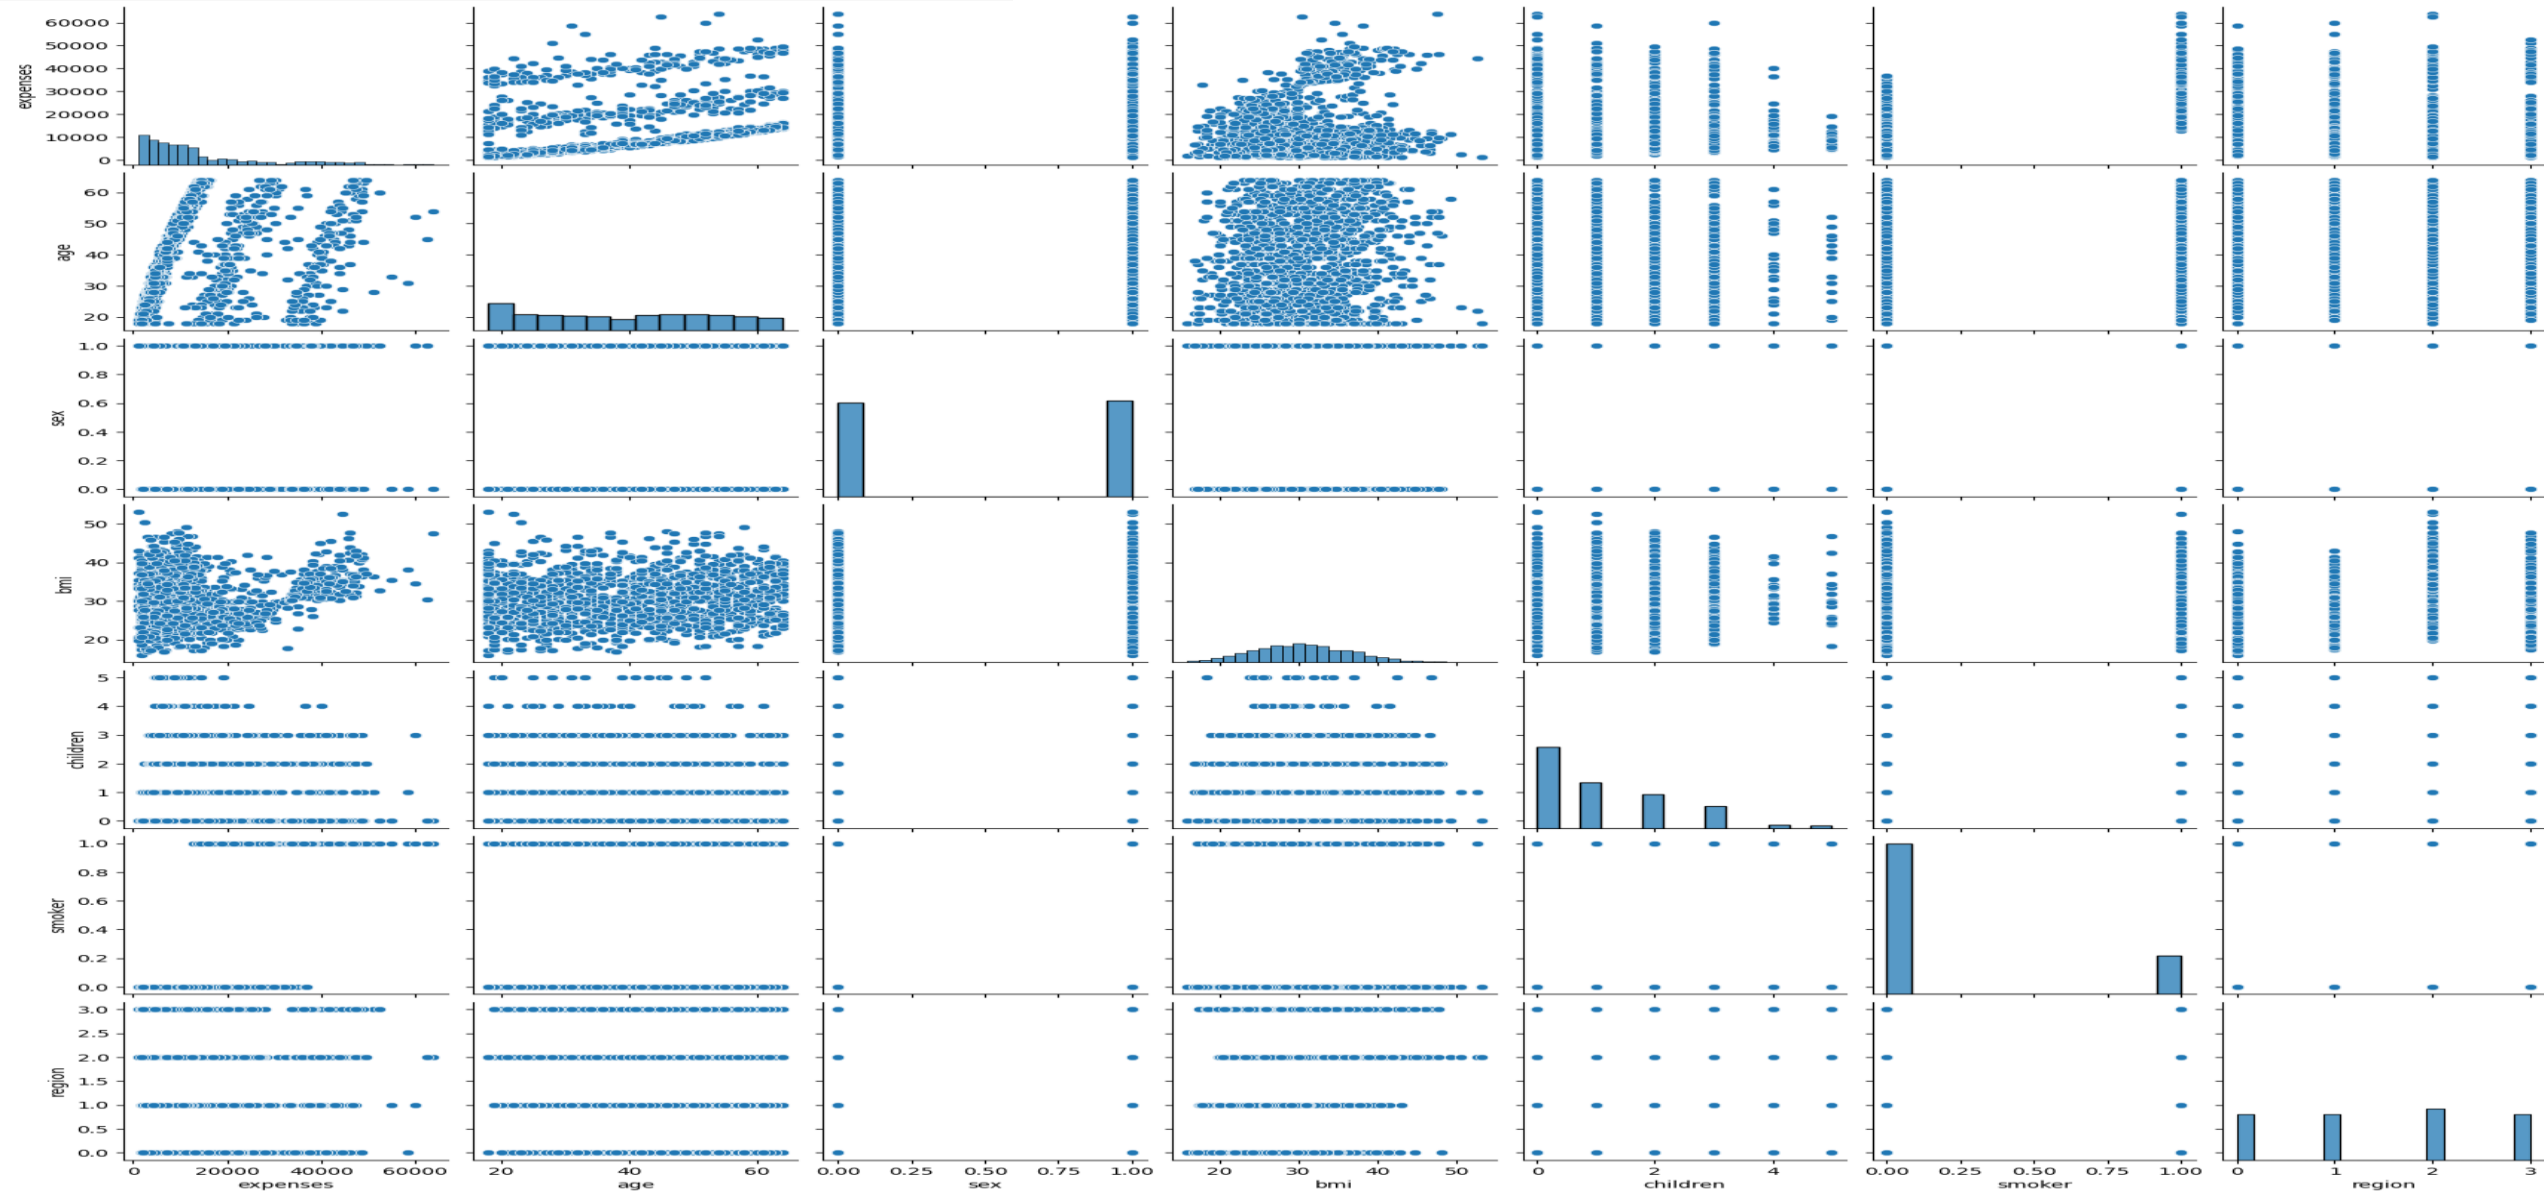
Linear-Regression

```
df_corr = df.corr()  
df_corr_sort = df_corr.sort_values("expenses", ascending=False)  
df_corr_sort["expenses"].head(10)
```

```
expenses    1.000000  
smoker      0.787251  
age         0.299008  
bmi         0.198576  
children    0.067998  
sex         0.057292  
region     -0.006208  
Name: expenses, dtype: float64
```


Linear-Regression

```
cols = ['expenses', 'age', 'sex', 'bmi', 'children', 'smoker', 'region']  
sns.pairplot(df[cols])  
plt.show()
```



Linear-Regression

```
colormap = plt.cm.gist_heat
plt.figure(figsize=(7,7))
sns.heatmap(df.corr(), linewidths=0.1, vmax=0.5, cmap=colormap, linecolor='white', annot=True)
plt.show()
```



Linear-Regression

```
cols_train = ['age', 'sex', 'bmi', 'children', 'smoker', 'region']  
X_train_pre = df[cols_train]  
y = df['expenses'].values  
  
X_train, X_test, y_train, y_test = train_test_split(X_train_pre, y, test_size=0.2, shuffle=True)
```

```
model = LinearRegression()  
model.fit(X_train, y_train)
```

```
y_pred = model.predict(X_test)
```

```
y_pred = model.predict(X_test)
```


```
# 성능 평가  
mae = mean_absolute_error(y_test, y_pred)  
mse = mean_squared_error(y_test, y_pred)  
rmse = np.sqrt(mse)  
r2 = r2_score(y_test, y_pred)
```

```
# 결과 출력  
print('Mean Absolute Error:', mae)  
print('Mean Squared Error:', mse)  
print('Root Mean Squared Error:', rmse)  
print('R-squared Score:', r2)
```



Mean Absolute Error: 3974.4165959075804
Mean Squared Error: 31658845.955824777
Root Mean Squared Error: 5626.619407408393
R-squared Score: 0.7461054175658259

Linear-Regression



```
# 예측 결과 시각화
plt.figure(figsize=(8, 6))
sns.scatterplot(x=y_test, y=y_pred)
plt.xlabel('Actual Expenses')
plt.ylabel('Predicted Expenses')
plt.title('Actual vs Predicted Expenses')
plt.show()
```

Linear-Regression

```
model = Sequential()

model.add(Dense(81, input_dim=6, activation='linear'))
model.add(Dense(1, activation='linear'))
model.summary()

modelpath = "./data/model/all{epoch:02d}-{val_accuracy:.4f}.hdf5"
checkpointer = ModelCheckpoint(filepath = modelpath, verbose=0, save_best_only = True)
early_stopping = EarlyStopping(monitor='val_loss', patience=10)

model.compile(loss='mean_squared_error', optimizer='adam', metrics=['accuracy'])
history = model.fit(X_train, y_train, epochs=2000, batch_size=500, validation_split=0.2, callbacks=[checkpointer, early_stopping])
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 81)	567
dense_7 (Dense)	(None, 1)	82
Total params: 649		
Trainable params: 649		
Non-trainable params: 0		

Linear-Regression

```
Epoch 1985/2000
2/2 [=====] - 0s 79ms/step - loss: 131905912.0000 - accuracy: 0.0000e+00 - val_loss: 126473120.0000 - val_accuracy: 0.0000e+00
Epoch 1986/2000
2/2 [=====] - 0s 82ms/step - loss: 131902800.0000 - accuracy: 0.0000e+00 - val_loss: 126470504.0000 - val_accuracy: 0.0000e+00
Epoch 1987/2000
2/2 [=====] - 0s 83ms/step - loss: 131899912.0000 - accuracy: 0.0000e+00 - val_loss: 126467752.0000 - val_accuracy: 0.0000e+00
Epoch 1988/2000
2/2 [=====] - 0s 61ms/step - loss: 131897000.0000 - accuracy: 0.0000e+00 - val_loss: 126465456.0000 - val_accuracy: 0.0000e+00
Epoch 1989/2000
2/2 [=====] - 0s 59ms/step - loss: 131893456.0000 - accuracy: 0.0000e+00 - val_loss: 126462728.0000 - val_accuracy: 0.0000e+00
Epoch 1990/2000
2/2 [=====] - 0s 77ms/step - loss: 131890536.0000 - accuracy: 0.0000e+00 - val_loss: 126460344.0000 - val_accuracy: 0.0000e+00
Epoch 1991/2000
2/2 [=====] - 0s 84ms/step - loss: 131887184.0000 - accuracy: 0.0000e+00 - val_loss: 126457696.0000 - val_accuracy: 0.0000e+00
Epoch 1992/2000
2/2 [=====] - 0s 77ms/step - loss: 131883984.0000 - accuracy: 0.0000e+00 - val_loss: 126454960.0000 - val_accuracy: 0.0000e+00
Epoch 1993/2000
2/2 [=====] - 0s 84ms/step - loss: 131881016.0000 - accuracy: 0.0000e+00 - val_loss: 126452304.0000 - val_accuracy: 0.0000e+00
Epoch 1994/2000
2/2 [=====] - 0s 79ms/step - loss: 131879904.0000 - accuracy: 0.0000e+00 - val_loss: 126449160.0000 - val_accuracy: 0.0000e+00
Epoch 1995/2000
2/2 [=====] - 0s 80ms/step - loss: 131875104.0000 - accuracy: 0.0000e+00 - val_loss: 126447336.0000 - val_accuracy: 0.0000e+00
Epoch 1996/2000
2/2 [=====] - 0s 62ms/step - loss: 131871504.0000 - accuracy: 0.0000e+00 - val_loss: 126445128.0000 - val_accuracy: 0.0000e+00
Epoch 1997/2000
2/2 [=====] - 0s 76ms/step - loss: 131868536.0000 - accuracy: 0.0000e+00 - val_loss: 126442920.0000 - val_accuracy: 0.0000e+00
Epoch 1998/2000
2/2 [=====] - 0s 75ms/step - loss: 131865344.0000 - accuracy: 0.0000e+00 - val_loss: 126440480.0000 - val_accuracy: 0.0000e+00
Epoch 1999/2000
2/2 [=====] - 0s 79ms/step - loss: 131862176.0000 - accuracy: 0.0000e+00 - val_loss: 126437760.0000 - val_accuracy: 0.0000e+00
Epoch 2000/2000
2/2 [=====] - 0s 63ms/step - loss: 131859848.0000 - accuracy: 0.0000e+00 - val_loss: 126435528.0000 - val_accuracy: 0.0000e+00
```

Linear-Regression

Linear-Regression은 loss값이 높게 나오는 것이 일반적인 현상이다.


손실 값이 높다고 모델이 잘못되었거나 문제가 있는 것은 아니다. 평균 제곱 오차를 사용하여 손실을 계산하며, 손실 값은 데이터의 크기와 범위에 따라 달라질 수 있다. 손실 값 자체보다는 R-squared score를 주로 사용하여 모델의 성능을 평가한다.

Linear-Regression

```
# Loss 값 추출
loss = history.history['loss']
val_loss = history.history['val_loss']

# 그래프 그리기
epochs = range(1, len(loss) + 1)

plt.plot(epochs, loss, 'r', label='train_loss')
plt.plot(epochs, val_loss, 'b', label='val_loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



Binary-CrossEntropy

수질 판별 후 식수 가능한지?

Binary-CrossEntropy

데이터셋 분석

pH – 물의 산성 or 알칼리성 정도

Hardness – 물의 칼슘과 마그네슘 함량

Solids – 물에 용해된 총 불순물의 양

Chloramines – 물 처리에 사용되는 염소와 암모니아의 혼합물(소독효과)

Sulfate – 물에 용해된 황산염의 양

Conductivity – 물의 전기 전도도

Organic_carbon – 물에 용해된 유기 탄소의 양

Trihalomethanes – 물 처리에 의해 생성된 삼할로메탄류 화합물의 양

Turbidity – 물의 탁도

Potability – 물의 음용 가능 여부

Binary-CrossEntropy

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from sklearn.model_selection import train_test_split
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

```
df = pd.read_csv('/content/drive/MyDrive/ai/water_potability.csv')
```

```
df.head(5)
df.isnull().sum().sort_values(ascending=False)
```



ph	0
Hardness	0
Solids	0
Chloramines	0
Sulfate	0
Conductivity	0
Organic_carbon	0
Trihalomethanes	0
Turbidity	0
Potability	0
dtype: int64	

Binary-CrossEntropy

df

	ph	Hardness	Solids	Chloramines	Sulfate	Conductivity	Organic_carbon	Trihalomethanes	Turbidity	Potability
0	8.316766	214.373394	22018.41744	8.059332	356.886136	363.266516	18.436525	100.341674	4.628771	0
1	9.092223	181.101509	17978.98634	6.546600	310.135738	398.410813	11.558279	31.997993	4.075075	0
2	5.584087	188.313324	28748.68774	7.544869	326.678363	280.467916	8.399735	54.917862	2.559708	0
3	10.223862	248.071735	28749.71654	7.513408	393.663395	283.651634	13.789695	84.603556	2.672989	0
4	8.635849	203.361523	13672.09176	4.563009	303.309771	474.607645	12.363817	62.798309	4.401425	0
...
2006	8.989900	215.047358	15921.41202	6.297312	312.931021	390.410231	9.899115	55.069304	4.613843	1
2007	6.702547	207.321086	17246.92035	7.708117	304.510230	329.266002	16.217303	28.878601	3.442983	1
2008	11.491011	94.812545	37188.82602	9.263166	258.930600	439.893618	16.172755	41.558501	4.369264	1
2009	6.069616	186.659040	26138.78019	7.747547	345.700257	415.886955	12.067620	60.419921	3.669712	1
2010	4.668102	193.681736	47580.99160	7.166639	359.948574	526.424171	13.894419	66.687695	4.435821	1

2011 rows × 10 columns

Binary-CrossEntropy

```
df["Potability"].value_counts()
```

```
0    1200
```

```
1     811
```

```
Name: Potability, dtype: int64
```

```
df_corr = df.corr()
```

```
df_corr_sort = df_corr.sort_values("Potability", ascending=False)
```

```
df_corr_sort['Potability'].head(10)
```

```
Potability    1.000000
```

```
Solids         0.040674
```

```
Turbidity      0.022682
```

```
Chloramines    0.020784
```

```
ph             0.014530
```

```
Trihalomethanes 0.009244
```

```
Hardness       -0.001505
```

```
Sulfate        -0.015303
```

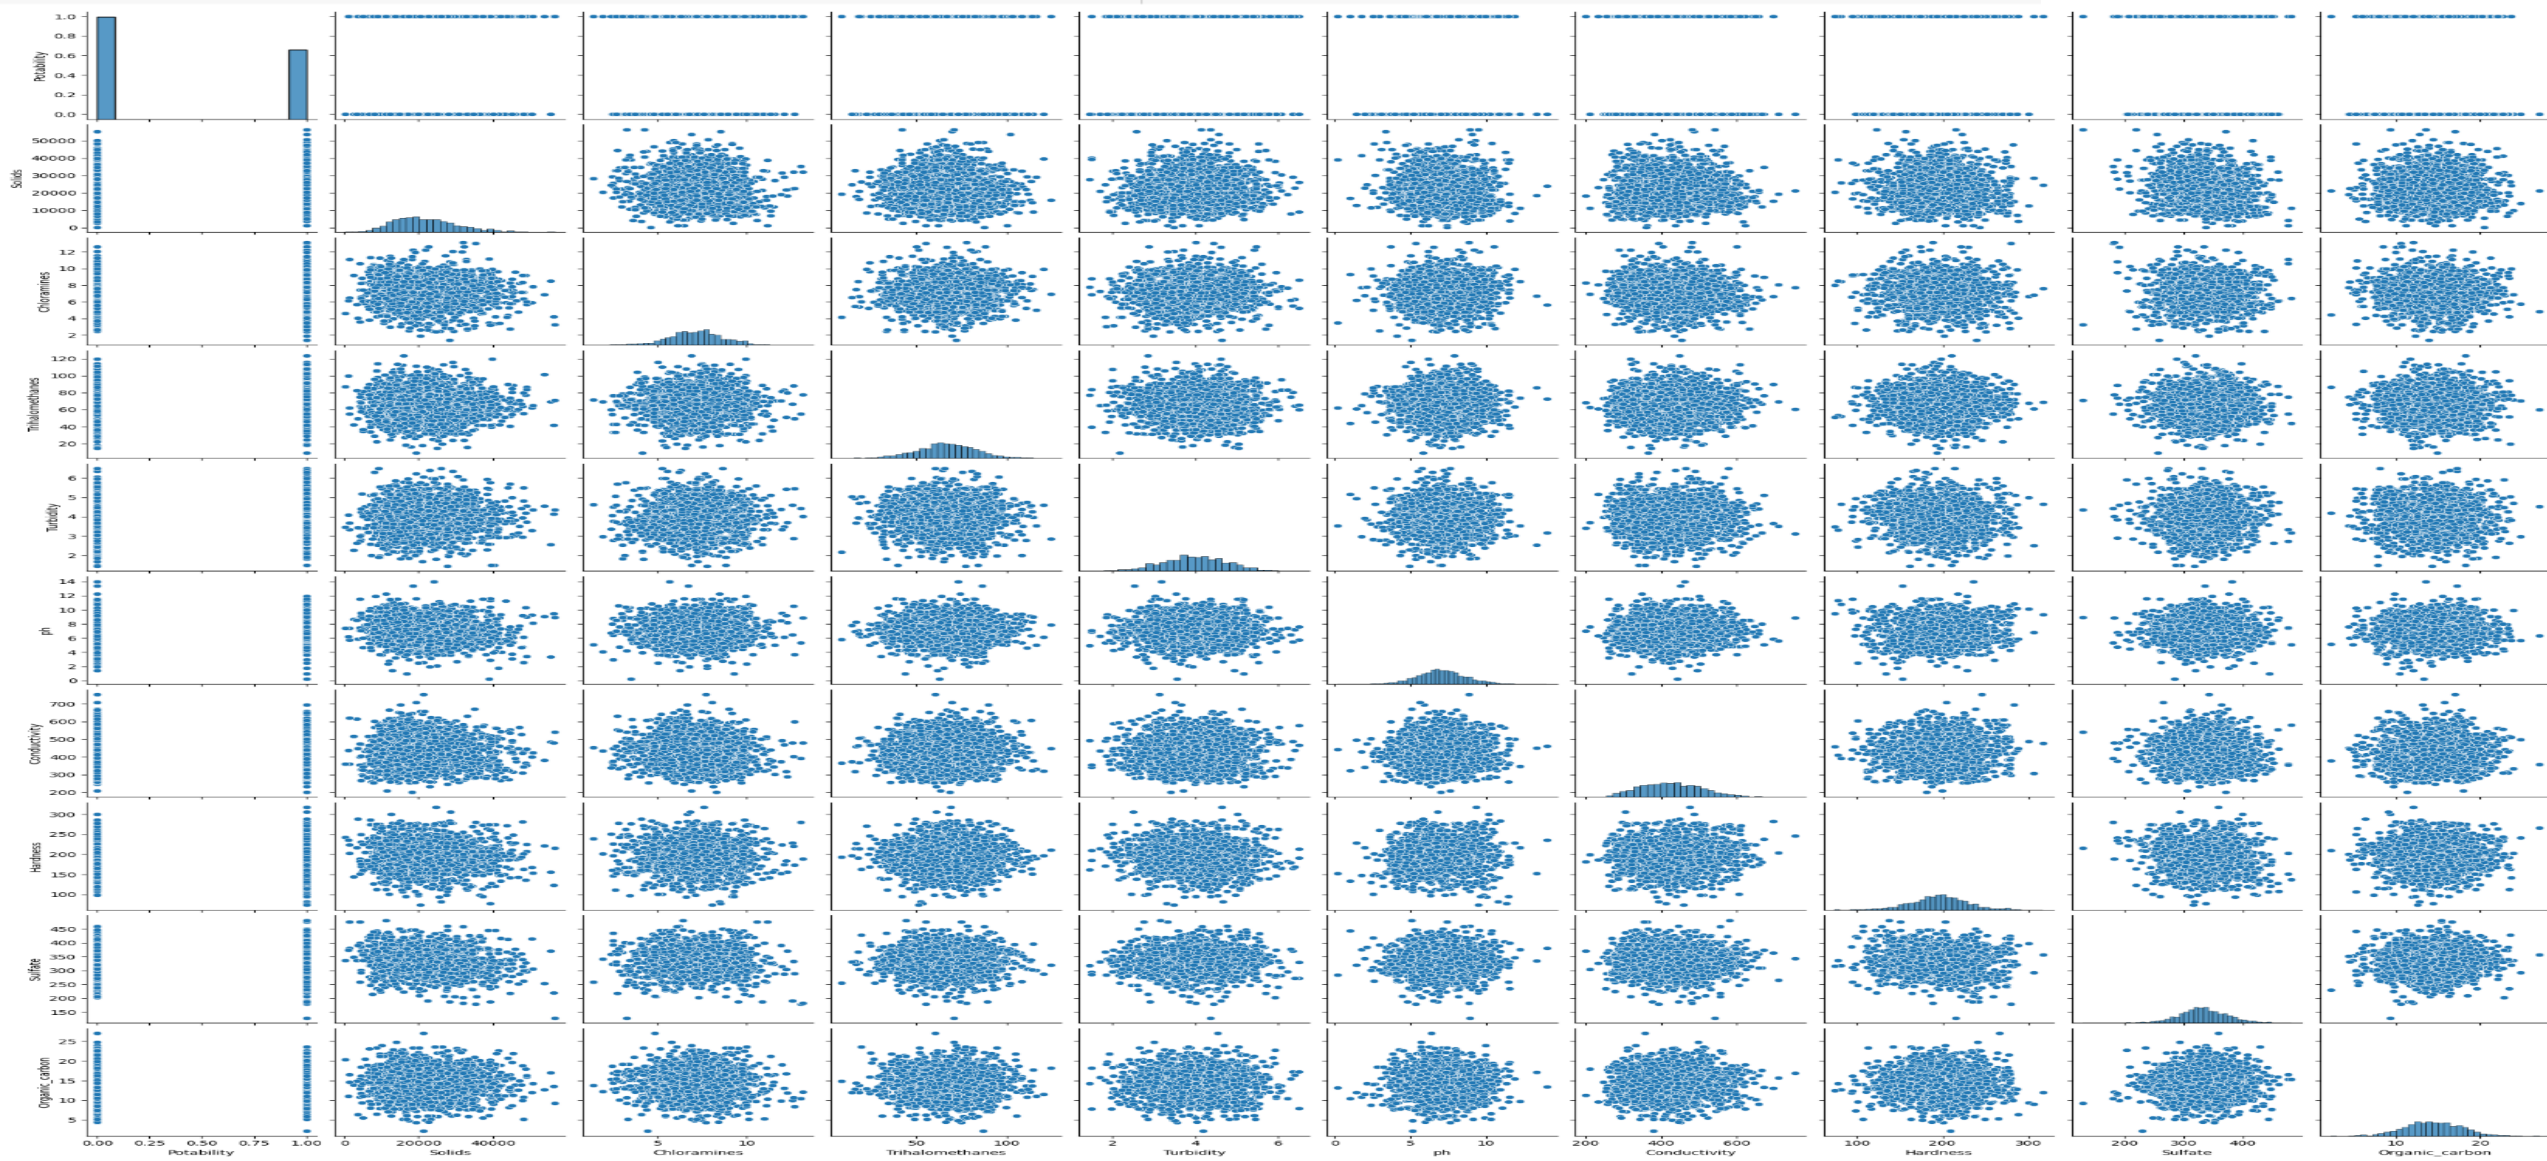
```
Conductivity   -0.015496
```

```
Organic_carbon -0.015567
```

```
Name: Potability, dtype: float64
```


Binary-CrossEntropy

```
cols = ['Potability', 'Solids', 'Chloramines', 'Trihalomethanes', 'Turbidity', 'ph', 'Conductivity', 'Hardness', 'Sulfate', 'Organic_carbon']  
sns.pairplot(df[cols])  
plt.show()
```



Binary-CrossEntropy

```
colormap = plt.cm.gist_heat
plt.figure(figsize=(12,12))
sns.heatmap(df.corr(), linewidths=0.1, vmax=0.5, cmap=colormap, linecolor='white', annot=True)
plt.show()
```



Binary-CrossEntropy

```
print(df.shape[1])
```

```
10
```

```
cols_train = ['Solids', 'Chloramines', 'Trihalomethanes', 'Turbidity', 'ph', 'Conductivity', 'Hardness', 'Sulfate', 'Organic_carbon' ]  
X_train_pre = df[cols_train]  
y = df['Potability'].values
```

```
X_train, X_test, y_train, y_test = train_test_split(X_train_pre, y, test_size=0.2, shuffle=True)
```

```
model = Sequential()
```

```
model.add(Dense(100, input_dim=9, activation='relu'))  
model.add(Dense(90, activation='relu'))  
model.add(Dense(100, activation='relu'))  
model.add(Dense(50, activation='relu'))  
model.add(Dense(100, activation='relu'))  
model.add(Dense(50, activation='relu'))  
model.add(Dense(25, activation='relu'))  
model.add(Dense(25, activation='relu'))  
model.add(Dense(5, activation='relu'))  
model.add(Dense(1, activation='sigmoid'))  
model.summary()
```

```
modelpath = "./data/model/all{epoch:02d}-{val_accuracy:.4f}.hdf5"
```

```
checkpointer = ModelCheckpoint(filepath = modelpath, verbose=0, save_best_only = True)
```

```
early_stopping = EarlyStopping(monitor='val_loss', patience=10)
```

```
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
history = model.fit(X_train, y, epochs=2000, batch_size=50, validation_split=0.2, callbacks=[checkpointer, early_stopping])
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_20 (Dense)	(None, 100)	1000
dense_21 (Dense)	(None, 90)	9090
dense_22 (Dense)	(None, 100)	9100
dense_23 (Dense)	(None, 50)	5050
dense_24 (Dense)	(None, 100)	5100
dense_25 (Dense)	(None, 50)	5050
dense_26 (Dense)	(None, 25)	1275
dense_27 (Dense)	(None, 25)	650
dense_28 (Dense)	(None, 5)	130
dense_29 (Dense)	(None, 1)	6

=====
Total params: 36,451
Trainable params: 36,451
Non-trainable params: 0

Binary-CrossEntropy


```
26/26 [=====] - 0s 4ms/step - loss: 0.6686 - accuracy: 0.6104 - val_loss: 0.6443 - val_accuracy: 0.6646
Epoch 77/2000
26/26 [=====] - 0s 4ms/step - loss: 0.6686 - accuracy: 0.6104 - val_loss: 0.6443 - val_accuracy: 0.6646
Epoch 78/2000
26/26 [=====] - 0s 5ms/step - loss: 0.6686 - accuracy: 0.6104 - val_loss: 0.6443 - val_accuracy: 0.6646
Epoch 79/2000
26/26 [=====] - 0s 5ms/step - loss: 0.6686 - accuracy: 0.6104 - val_loss: 0.6443 - val_accuracy: 0.6646
Epoch 80/2000
26/26 [=====] - 0s 5ms/step - loss: 0.6686 - accuracy: 0.6104 - val_loss: 0.6443 - val_accuracy: 0.6646
Epoch 81/2000
26/26 [=====] - 0s 6ms/step - loss: 0.6686 - accuracy: 0.6104 - val_loss: 0.6442 - val_accuracy: 0.6646
Epoch 82/2000
26/26 [=====] - 0s 4ms/step - loss: 0.6686 - accuracy: 0.6104 - val_loss: 0.6442 - val_accuracy: 0.6646
Epoch 83/2000
26/26 [=====] - 0s 5ms/step - loss: 0.6686 - accuracy: 0.6104 - val_loss: 0.6442 - val_accuracy: 0.6646
Epoch 84/2000
26/26 [=====] - 0s 4ms/step - loss: 0.6686 - accuracy: 0.6104 - val_loss: 0.6442 - val_accuracy: 0.6646
Epoch 85/2000
26/26 [=====] - 0s 5ms/step - loss: 0.6686 - accuracy: 0.6104 - val_loss: 0.6443 - val_accuracy: 0.6646
Epoch 86/2000
26/26 [=====] - 0s 5ms/step - loss: 0.6686 - accuracy: 0.6104 - val_loss: 0.6443 - val_accuracy: 0.6646
Epoch 87/2000
26/26 [=====] - 0s 5ms/step - loss: 0.6686 - accuracy: 0.6104 - val_loss: 0.6443 - val_accuracy: 0.6646
Epoch 88/2000
26/26 [=====] - 0s 5ms/step - loss: 0.6686 - accuracy: 0.6104 - val_loss: 0.6442 - val_accuracy: 0.6646
Epoch 89/2000
26/26 [=====] - 0s 5ms/step - loss: 0.6686 - accuracy: 0.6104 - val_loss: 0.6442 - val_accuracy: 0.6646
Epoch 90/2000
26/26 [=====] - 0s 5ms/step - loss: 0.6686 - accuracy: 0.6104 - val_loss: 0.6442 - val_accuracy: 0.6646
Epoch 91/2000
26/26 [=====] - 0s 4ms/step - loss: 0.6686 - accuracy: 0.6104 - val_loss: 0.6443 - val_accuracy: 0.6646
```


Binary-CrossEntropy

```
# Loss 값 추출
loss = history.history['loss']
val_loss = history.history['val_loss']

# 그래프 그리기
epochs = range(1, len(loss) + 1)

plt.plot(epochs, loss, 'r', label='train_loss')
plt.plot(epochs, val_loss, 'b', label='val_loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



Multi-classfication

와인 품질 예측

Multi-classfication

데이터 셋 분석

- 1 - fixed acidity
- 2 - volatile acidity
- 3 - citric acid
- 4 - residual sugar
- 5 - chlorides
- 6 - free sulfur dioxide
- 7 - total sulfur dioxide
- 8 - density
- 9 - pH
- 10 - sulphates
- 11 - alcohol
- 12 - quality (score between 0 and 10)

1. 고정산도
2. 휘발성산도
3. 구연산
4. 잔류 당분
5. 염화물
6. 유리황산
7. 총황산
8. 밀도
9. pH
10. 황산염
11. 알코올
12. 품질 (0에서 10까지의 점수)

Multi-classfication

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from sklearn.model_selection import train_test_split
from keras.utils import to_categorical
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

df = pd.read_csv('/content/drive/MyDrive/ai/WineQT.csv')
```

[+ 코드](#)[+ 텍스트](#)

```
df.head(5)
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	Id	quality
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	0	5
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	1	5
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	2	5
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	3	6
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	4	5

Multi-classfication

```
df.isnull().sum().sort_values(ascending=False)
```

```
fixed acidity      0  
volatile acidity  0  
citric acid       0  
residual sugar    0  
chlorides         0  
free sulfur dioxide 0  
total sulfur dioxide 0  
density          0  
pH               0  
sulphates        0  
alcohol          0  
Id               0  
quality          0  
dtype: int64
```

```
df["quality"].value_counts()  
df
```



	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	Id	quality
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	0	5
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8	1	5
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8	2	5
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8	3	6
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	4	5
...
1138	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11.0	1592	6
1139	6.8	0.620	0.08	1.9	0.068	28.0	38.0	0.99651	3.42	0.82	9.5	1593	6
1140	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5	1594	5
1141	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11.2	1595	6
1142	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10.2	1597	5

1143 rows x 13 columns

Multi-classfication

```
) df_corr = df.corr()  
df_corr_sort = df_corr.sort_values("quality", ascending=False)  
df_corr_sort['quality'].head(12)
```

quality	1.000000
alcohol	0.484866
sulphates	0.257710
citric acid	0.240821
fixed acidity	0.121970
ld	0.069708
residual sugar	0.022002
pH	-0.052453
free sulfur dioxide	-0.063260
chlorides	-0.124085
density	-0.175208
total sulfur dioxide	-0.183339

Name: quality, dtype: float64

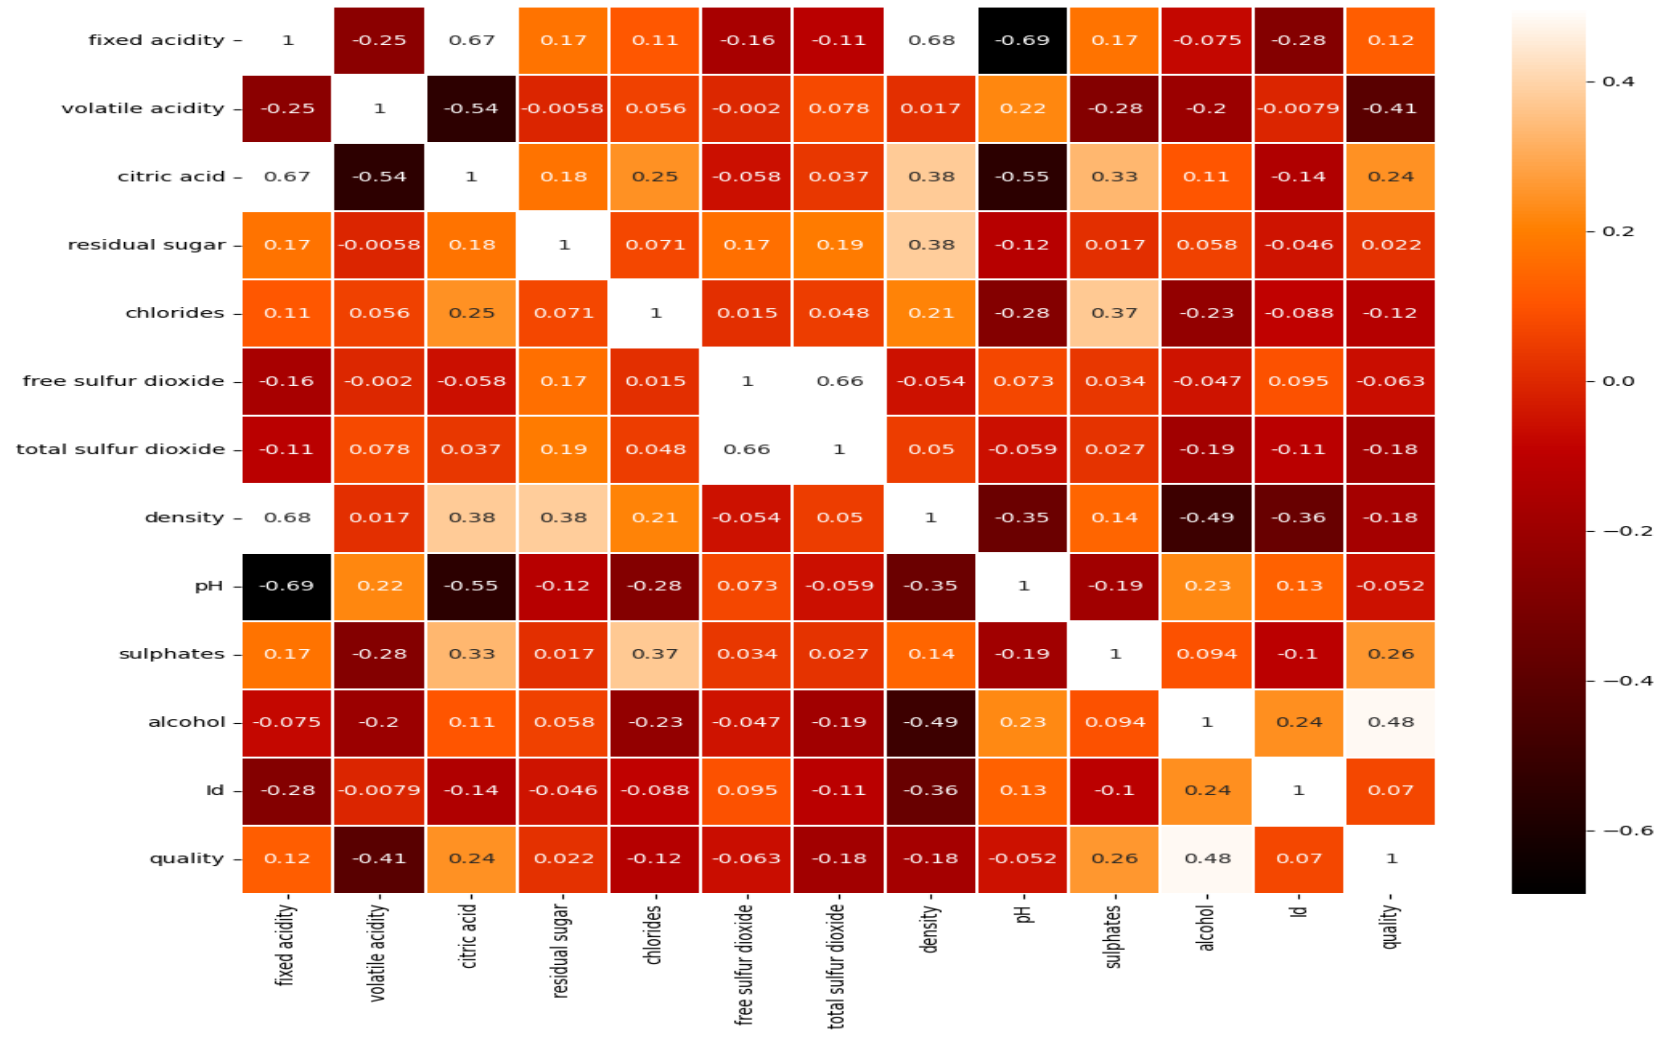
Multi-classfication

```
cols = ['quality', 'alcohol', 'sulphates', 'citric acid', 'fixed acidity', 'residual sugar', 'pH', 'free sulfur dioxide', 'chlorides', 'density']  
sns.pairplot(df[cols])  
plt.show()
```



Multi-classfication

```
colormap = plt.cm.gist_heat
plt.figure(figsize=(12,12))
sns.heatmap(df.corr(), linewidths=0.1, vmax=0.5, cmap=colormap, linecolor='white', annot=True)
plt.show()
```



Multi-classfication

```
print(df.shape[1])
```

```
13
```

```
cols_train = ["fixed acidity", "volatile acidity", "citric acid", "residual sugar", "chlorides", "free sulfur dioxide",  
              "total sulfur dioxide", "density", "pH", "sulphates", "alcohol", "id"]
```

```
X_train_pre = df[cols_train]
```

```
y = df['quality'].values
```

```
num_classes = 9
```

```
X_train, X_test, y_train, y_test = train_test_split(X_train_pre, y, test_size=0.2, shuffle=True)
```

```
y_train_encoded = to_categorical(y_train, num_classes)
```

```
model = Sequential()
```

```
num_classes = 9
```

```
model.add(Dense(72, input_dim=12, activation='relu'))
```

```
model.add(Dense(144, activation='relu'))
```

```
model.add(Dense(144, activation='relu'))
```

```
model.add(Dense(144, activation='relu'))
```

```
model.add(Dense(72, activation='relu'))
```

```
model.add(Dense(num_classes, activation='softmax'))
```

```
modelpath = "./data/model/all{epoch:02d}-{val_accuracy:.4f}.hdf5"
```

```
checkpointer = ModelCheckpoint(filepath = modelpath, verbose=0, save_best_only = True)
```

```
early_stopping = EarlyStopping(monitor='val_loss', patience=50)
```

```
model.summary()
```

```
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
history = model.fit(X_train, y_train_encoded, epochs=2000, batch_size=300, validation_split=0.2, callbacks=[checkpointer, early_stopping])
```

Multi-classfication

Model: "sequential_31"

Layer (type)	Output Shape	Param #	Epoch 81/2000 3/3 [=====] - 0s 16ms/step - loss: 0.2085 - accuracy: 0.5732 - val_loss: 0.2839 - val_accuracy: 0.4262
dense_188 (Dense)	(None, 72)	936	Epoch 82/2000 3/3 [=====] - 0s 25ms/step - loss: 0.2119 - accuracy: 0.5732 - val_loss: 0.2472 - val_accuracy: 0.4317
dense_189 (Dense)	(None, 144)	10512	Epoch 83/2000 3/3 [=====] - 0s 14ms/step - loss: 0.2098 - accuracy: 0.5321 - val_loss: 0.2905 - val_accuracy: 0.3934
dense_190 (Dense)	(None, 144)	20880	Epoch 84/2000 3/3 [=====] - 0s 17ms/step - loss: 0.2214 - accuracy: 0.5103 - val_loss: 0.2619 - val_accuracy: 0.4044
dense_191 (Dense)	(None, 144)	20880	Epoch 85/2000 3/3 [=====] - 0s 19ms/step - loss: 0.2053 - accuracy: 0.5486 - val_loss: 0.2367 - val_accuracy: 0.4426
dense_192 (Dense)	(None, 72)	10440	Epoch 86/2000 3/3 [=====] - 0s 17ms/step - loss: 0.2027 - accuracy: 0.5267 - val_loss: 0.2513 - val_accuracy: 0.4262
dense_193 (Dense)	(None, 9)	657	Epoch 87/2000 3/3 [=====] - 0s 17ms/step - loss: 0.2023 - accuracy: 0.5540 - val_loss: 0.2457 - val_accuracy: 0.4262
Total params: 64,305			Epoch 88/2000 3/3 [=====] - 0s 17ms/step - loss: 0.1988 - accuracy: 0.5691 - val_loss: 0.2428 - val_accuracy: 0.4153
Trainable params: 64,305			Epoch 89/2000 3/3 [=====] - 0s 21ms/step - loss: 0.1998 - accuracy: 0.5595 - val_loss: 0.2436 - val_accuracy: 0.4262
Non-trainable params: 0			Epoch 90/2000 3/3 [=====] - 0s 21ms/step - loss: 0.1973 - accuracy: 0.5650 - val_loss: 0.2380 - val_accuracy: 0.4481
			Epoch 91/2000 3/3 [=====] - 0s 19ms/step - loss: 0.2087 - accuracy: 0.5198 - val_loss: 0.2365 - val_accuracy: 0.4262
			Epoch 92/2000 3/3 [=====] - 0s 19ms/step - loss: 0.2056 - accuracy: 0.5075 - val_loss: 0.2440 - val_accuracy: 0.4153
			Epoch 93/2000 3/3 [=====] - 0s 20ms/step - loss: 0.2030 - accuracy: 0.5376 - val_loss: 0.2776 - val_accuracy: 0.3880
			Epoch 94/2000 3/3 [=====] - 0s 17ms/step - loss: 0.2077 - accuracy: 0.5540 - val_loss: 0.2347 - val_accuracy: 0.4372
			Epoch 95/2000 3/3 [=====] - 0s 18ms/step - loss: 0.2033 - accuracy: 0.5636 - val_loss: 0.2464 - val_accuracy: 0.4317
			Epoch 96/2000 3/3 [=====] - 0s 18ms/step - loss: 0.2011 - accuracy: 0.5554 - val_loss: 0.2385 - val_accuracy: 0.4153
			Epoch 97/2000 3/3 [=====] - 0s 17ms/step - loss: 0.2003 - accuracy: 0.5431 - val_loss: 0.2454 - val_accuracy: 0.4098
			Epoch 98/2000 3/3 [=====] - 0s 18ms/step - loss: 0.2021 - accuracy: 0.5445 - val_loss: 0.2983 - val_accuracy: 0.3880
			Epoch 99/2000 3/3 [=====] - 0s 18ms/step - loss: 0.2151 - accuracy: 0.5349 - val_loss: 0.2463 - val_accuracy: 0.4262
			Epoch 100/2000 3/3 [=====] - 0s 18ms/step - loss: 0.2113 - accuracy: 0.5417 - val_loss: 0.2609 - val_accuracy: 0.4372

Multi-classfication

```
# Loss 값 추출
loss = history.history['loss']
val_loss = history.history['val_loss']

# 그래프 그리기
epochs = range(1, len(loss) + 1)

plt.plot(epochs, loss, 'r', label='train_loss')
plt.plot(epochs, val_loss, 'b', label='val_loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

